

VE414 Lecture 20

Jing Liu

UM-SJTU Joint Institute

July 25, 2019

- Consider the following problem

Bayesian Network for Text Mining

Suppose we have been given a collection of m documents and we wish to determine how the documents are related, and whether they can be sorted into k categories. i.e. in terms of topic, author or era in which they were written.

- For example, consider a two-topic model of news, with one on “politics” and the other on “business.” The most common words for politics might be

“President”, “Congress”, and “government”,

while the business topic may be made up of words such as

“profit”, “budget”, and “market”.

- However, words can be equally important for both topics, i.e. a word like
“expect”.

Q: What is a sensible Bayesian Network model to use?





Plate Diagram for LDA



$$\phi_k \sim \text{Dirichlet}(\beta)$$

$$\theta_i \sim \text{Dirichlet}(\alpha)$$

$$z_{i,j} \sim \text{Multinomial}(\theta_i)$$

$$w_{i,j} \sim \text{Multinomial}(\phi_{z_{i,j}})$$


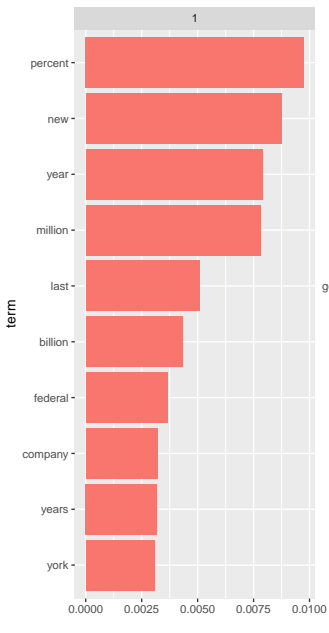
Q: What is the joint posterior distribution?


- beta's here are components of ϕ_k

```
> ap_lda = LDA(AssociatedPress, k = 2,
+             method = "Gibbs",
+             control = list(seed = 1234))
>
> ap_topics = tidy(ap_lda, matrix = "beta")
> ap_topics
```

```
# A tibble: 20,946 x 3
  topic term      beta
  <int> <chr>    <dbl>
1     1 aaron    0.000000501
2     2 aaron    0.0000424
3     1 abandon  0.000000501
4     2 abandon  0.0000634
5     1 abandoned 0.0000255
6     2 abandoned 0.000147
7     1 abandoning 0.000000501
8     2 abandoning 0.0000256
9     1 abbott    0.0000406
10    2 abbott    0.000000420
# with 20,936 more rows
```

- It is often interesting to see terms that are most common within each topic.





- gamma's here are components of θ_i

```
> ap_documents = tidy(ap_lda, matrix = "gamma")  
> ap_documents
```

```
# A tibble: 4,492 x 3  
  document topic gamma  
  <int> <int> <dbl>  
1         1     1 0.192  
2         2     1 0.411  
3         3     1 0.272  
4         4     1 0.405  
5         5     1 0.457  
6         6     1 0.240  
7         7     1 0.641  
8         8     1 0.362  
9         9     1 0.526  
10        10     1 0.301  
# with 4,482 more rows
```

```
> ap_documents[ap_documents$document==1,]
```

```
# A tibble: 2 x 3  
  document topic gamma  
  <int> <int> <dbl>  
1         1     1 0.192  
2         1     2 0.808
```



```
> tidy(AssociatedPress) %>%  
+   filter(document == 1) %>%  
+   arrange(desc(count))
```

```
# A tibble: 186 x 3  
  document term      count  
  <int> <chr>    <dbl>  
1       1 police      7  
2       1 school      7  
3       1 teacher      7  
4       1 shot        5  
5       1 students     5  
6       1 boy         4  
7       1 boys        4  
8       1 classroom    4  
9       1 gun         3  
10      1 guns         3  
11      1 jammed       3  
12      1 marino       3  
13      1 yearold      3  
14      1 adult        2  
15      1 apparently    2  
# with 171 more rows
```

- This looks like a political news instead of business news, which means LDA has correctly sorted the document in this case.

Q: Anyone uses twitter? Anyone follow Donald Trump's twitter account?

- There was a claim about his tweets circulating in 2016:

Todd Vaziri

Every non-hyperbolic tweet is from iPhone.

i.e. his staff handled

Every hyperbolic tweet is from Android.

i.e. really from him



```
> library(twitterR)
> # access to the twitter API.
> consumer_key = "your_consumer_key"
> consumer_secret = "your_consumer_secret"
> access_token = "your_access_token"
> access_secret = "your_access_secret"
> setup_twitter_oauth(consumer_key, consumer_secret,
+                      access_token, access_secret)
```

```
> library(dplyr)
>
> trump_tweets =
+   userTimeline("realDonaldTrump", n = 5)
> class(trump_tweets)
```

```
[1] "list"
```

```
> (trump_tweets_tb =
+   as_tibble(
+     purrr::map_dfr(trump_tweets, as.data.frame)))
```

```
# A tibble: 5 x 16
  text          favorited favoriteCount replyToSN created          truncated
  <chr>          <lgl>          <dbl> <lgl>      <dtm>          <lgl>
1 Secretary Po? FALSE          53953. NA      2018-05-09 12:35:51 TRUE
2 I am pleased? FALSE          77505. NA      2018-05-09 12:30:56 TRUE
3 Congratulati? FALSE          32838. NA      2018-05-09 12:00:30 TRUE
4 Candace Owen? FALSE          62859. NA      2018-05-09 11:48:17 TRUE
5 The Fake New? FALSE          56652. NA      2018-05-09 11:38:45 TRUE
# ... with 10 more variables: replyToSID <lgl>, id <chr>, replyToUID <lgl>,
#   statusSource <chr>, screenName <chr>, retweetCount <dbl>, isRetweet <lgl>,
#   retweeted <lgl>, longitude <lgl>, latitude <lgl>
```

```

> rm(list = ls())
> # Tweets in 2016
> load("~/Desktop/trump_tweets.rda")
> # load R object that was saved
> trump_tweets_tb

```

```

# A tibble: 1,512 x 16
  text          favorited favoriteCount replyToSN created          truncated
  <chr>         <lgl>          <dbl> <chr>      <dtm>          <lgl>
1 My economic ? FALSE          9214. NA      2016-08-08 15:20:44 FALSE
2 Join me in F? FALSE          6981. NA      2016-08-08 13:28:20 FALSE
# ... with 1,510 more rows, and 10 more variables: replyToSID <lgl>, id <chr>,
#   replyToUID <chr>, statusSource <chr>, screenName <chr>, retweetCount <dbl>,
#   isRetweet <lgl>, retweeted <lgl>, longitude <chr>, latitude <chr>

```

```

> (sel_tb =
+   select(trump_tweets_tb,
+         id, statusSource, text, created))

```

```

# A tibble: 1,512 x 4
  id          statusSource      text          created
  <chr>        <chr>          <chr>      <dtm>
1 762669882571980801 "<a href=\"http://tw? My economic pol? 2016-08-08 15:20:44
2 762641595439190016 "<a href=\"http://tw? Join me in Faye? 2016-08-08 13:28:20
# ... with 1,510 more rows

```

```
> head(sel_tb$statusSource)
```

```
[1] "<a href=\".../android\" rel=\"nofollow\">Twitter for Android</a>"
[2] "<a href=\".../iphone\" rel=\"nofollow\">Twitter for iPhone</a>"
[3] "<a href=\".../iphone\" rel=\"nofollow\">Twitter for iPhone</a>"
[4] "<a href=\".../android\" rel=\"nofollow\">Twitter for Android</a>"
[5] "<a href=\".../android\" rel=\"nofollow\">Twitter for Android</a>"
[6] "<a href=\".../android\" rel=\"nofollow\">Twitter for Android</a>"
```

```
> (ext_tb =
+   extract(sel_tb,
+           col = statusSource,
+           into = "source",
+           regex = "Twitter for (.*)<"))
```

```
# A tibble: 1,512 x 4
   id          source text                                created
<chr>         <chr> <chr>                                <dtm>
1 762669882571980801 Android My economic policy speech wil? 2016-08-08 15:20:44
2 762641595439190016 iPhone  Join me in Fayetteville, Nort? 2016-08-08 13:28:20
# ... with 1,510 more rows
```

```
> unique(ext_tb$source)
```

```
[1] "Android" "iPhone"  NA        "iPad"
```

```

> trump_tidy_tb =
+   filter(ext_tb,
+           source %in% c("iPhone", "Android"))
>
> by_source = group_by(trump_tidy_tb, source)
> summarise(by_source, freq = n())

```

```

# A tibble: 2 x 2
  source      freq
  <chr>      <int>
1 Android      762
2 iPhone       628

```


- In practice, you would do those steps in one chunk using a compact syntax

```

> trump_tidy_tb = trump_tweets_tb %>%
+   select(id, statusSource, text, created) %>%
+   extract(statusSource, "source",
+           "Twitter for (.*)<") %>%
+   filter(source %in% c("iPhone", "Android"))

```

- Investigating the relation between % of tweets by source and time, we have



```

> trump_tidy_tb %>%
+   count(source, hour =
+     lubridate::hour(
+       lubridate::with_tz(
+         created, "EST")))) %>%
+   mutate(percent = n / sum(n)) %>%
+   ggplot(
+     aes(hour,
+         percent, color = source)
+   ) +
+   geom_line() +
+   scale_y_continuous(
+     labels = scales::percent_format()
+   ) +
+   labs(x = "Hour of day (EST)",
+        y = "% of tweets",
+        color = "")

```



- Another place we can spot a difference is in Trump's tendency of

“manually retweeting”


others by copying-pasting, then putting them under quotation marks.

```
> library(stringr) # Better handling on strings
> tweet_quotation_counts = trump_tidy_tb %>%
+   count(source, quotation =
+       ifelse(str_detect(text, '^"'),
+             "Manual Quotation",
+             "Using Retweeting Button"))
>
> ggplot(tweet_quotation_counts,
+       aes(source, n, fill = quotation)) +
+   geom_bar(stat = "identity",
+           position = "dodge") +
+   labs(x = "", y = "Number of tweets", fill = "")
```

- Almost all of those quoted tweets are posted from the android device.



- Another difference involves sharing links or pictures in tweets.



```

> tweet_picture_counts = trump_tidy_tb %>%
+   filter(!str_detect(text, '^\"')) %>%
+   # we have to remove retweeting cases
+   # that were done manually
+   count(source, picture =
+     ifelse(str_detect(text, "t.co"),
+       "Picture/link", "No picture/link")
+   )
> # twitter uses the domain https://t.co/
> # for all pictures and links, e.g.
> trump_tidy_tb$text[2]

```

```
[1] "Join me in Fayetteville, North Carolina tomorrow evening at 6pm. Tickets now
available at: https://t.co/Z80d4MYIg8"
```

```

> ggplot(tweet_picture_counts,
+   aes(source, n, fill = picture)) +
+   geom_bar(stat = "identity",
+     position = "dodge") +
+   labs(x = "", y = "Number of tweets", fill = "")

```

Q: What were the most common words in Trump's tweets overall?

```
> library(tidytext) # Good for tidy up strings
> reg = "([A-Za-z\\d#@']|'(?![A-Za-z\\d#@]))"
> # Separator between words in his tweets

> trump_words = trump_tidy_tb %>%
+   filter(!str_detect(text, '^"')) %>%
+   mutate(text = str_replace_all(
+     text, "https://t.co/[A-Za-z\\d]+|&";",
+     "")) %>%
+   # remove all pictures and links
+   unnest_tokens(word, text,
+     token = "regex",
+     pattern = reg) %>%
+   # split sentences into words
+   filter(!word %in% stop_words$word,
+     str_detect(word, "[a-z]"))
> # keep only relevant words
```

```
> trump_words
```


```
# A tibble: 8,753 x 4
  id          source created      word
<chr>      <chr>   <dtm>    <chr>
1 676494179216805888 iPhone 2015-12-14 20:09:15 record
2 676494179216805888 iPhone 2015-12-14 20:09:15 health
# ... with 8,751 more rows
```

```
> trump_tidy_tb
```

```
# A tibble: 1,390 x 4
  id          source text                created
<chr>      <chr>   <chr>                <dtm>
1 762669882571980801 Android My economic policy speech wil? 2016-08-08 15:20:44
2 762641595439190016 iPhone  Join me in Fayetteville, Nort? 2016-08-08 13:28:20
# ... with 1,388 more rows
```

```
> trump_words %>%
+   count(word, sort = TRUE) %>%
+   head(20) %>%
+   mutate(word = reorder(word, n)) %>%
+   ggplot(aes(word, n)) +
+   geom_bar(stat = "identity") +
+   ylab("Occurrences") +
+   coord_flip()
```

- Recall the data is for 2016, so no surprise there!



- We sort words according whether it is more likely to coming from an android

$$\log \left(\frac{\frac{\# \text{ Android} + 1}{\text{Total Android} + 1}}{\frac{\# \text{ iPhone} + 1}{\text{Total iPhone} + 1}} \right)$$

```
> android_iphone_ratios = trump_words %>%  
+   count(word, source) %>%  
+   # count the occurrence of a word by source  
+   spread(source, n, fill = 0) %>%  
+   # convert source into two columns  
+   mutate_at(c("Android", "iPhone"),  
+             funs((. + 1) / sum(. + 1))) %>%  
+   # apply a function two both columns  
+   mutate(logratio = log2(Android / iPhone)) %>%  
+   # create a new column  
+   arrange(desc(logratio))
```




```

> android_iphone_ratios %>%
+   group_by(logratio > 0) %>%
+   top_n(15, abs(logratio)) %>%
+   # 15 positive and 15 negative
+   ungroup() %>%
+   mutate(word = reorder(word, logratio)) %>%
+   # sort word according to logratio
+   ggplot(aes(word,
+               logratio,
+               fill = logratio < 0)) +
+   geom_bar(stat = "identity") +
+   coord_flip() +
+   ylab("Android / iPhone log ratio") +
+   scale_fill_manual(
+     name = "", labels = c("Android", "iPhone"),
+     values = c("red", "blue"))

```

Q: What can we say based on the following plot?



- The NRC emotion lexicon, which comes with `library(tidytext)`,

```
> (nrc = sentiments %>%  
+   filter(lexicon == "nrc") %>%  
+   select(word, sentiment))
```

```
# A tibble: 13,901 x 2  
  word      sentiment  
  <chr>    <chr>  
1 abacus    trust  
2 abandon   fear  
3 abandon   negative  
4 abandon   sadness  
# ... with 1.39e+04 more rows
```

is a way to associate common English words to 2 sentiments:

negative or positive

and 8 emotions:

anger, fear, anticipation, trust, surprise, sadness, joy, and disgust

- To measure the sentiment of the Android and iPhone tweets,

```
> trump_words
```

```
# A tibble: 8,753 x 4
  id          source created      word
<chr>      <chr>   <dtm>   <chr>
1 676494179216805888 iPhone 2015-12-14 20:09:15 record
2 676494179216805888 iPhone 2015-12-14 20:09:15 health
3 676494179216805888 iPhone 2015-12-14 20:09:15 #makeamericagreatagain
# ... with 8,750 more rows
```

we divide the number of Trump's words into each of the following categories

```
> unique(nrc$sentiment)
```

[1] "trust"	"fear"	"negative"	"sadness"	"anger"
[6] "surprise"	"positive"	"disgust"	"joy"	"anticipation"

```
> (join_tb = inner_join( # Join the two data sets
+   trump_words, nrc, by = "word"))
```

```
# A tibble: 5,109 x 5
  id          source created      word sentiment
<chr>      <chr>   <dtm>   <chr>   <chr>
1 676509769562251264 iPhone 2015-12-14 21:11:12 accolade anticipation
2 676509769562251264 iPhone 2015-12-14 21:11:12 accolade joy
3 676509769562251264 iPhone 2015-12-14 21:11:12 accolade positive
# ... with 5,106 more rows
```

```
> # Count the number of sentiment grouped by tweet
> (count_tb = count(join_tb, sentiment, id))
```

```
# A tibble: 3,722 x 3
  sentiment id          n
  <chr>      <chr>    <int>
1 anger    680503951440121856 1
2 anger    680734915718176768 1
3 anger    685490467329425408 1
# ... with 3,719 more rows
```

```
> # Add all possible combination of id and sentiment
> (complete_tb =
+   complete(count_tb,
+             sentiment, id, fill = list(n = 0)))
```

```
# A tibble: 8,790 x 3
  sentiment id          n
  <chr>      <chr>    <dbl>
1 anger    676509769562251264 0.
2 anger    680496083072593920 0.
3 anger    680503951440121856 1.
# ... with 8,787 more rows
```

- This dataset gives the counts of the 10 categories for each tweet.

```

> # Create a data set on tweets by source
> # One row for each of his tweets
> (sources = trump_words %>%
+   group_by(source) %>%
+   mutate(total = n()) %>%
+   # create a new variable
+   # total number of iPhone/Android
+   ungroup() %>%
+   distinct(id, source, total))

```

```

# A tibble: 1,172 x 3
  id          source total
<chr>      <chr>   <int>
1 676494179216805888 iPhone    3852
2 676509769562251264 iPhone    3852
3 680496083072593920 Android    4901
# ... with 1,169 more rows

```

```

> length(unique(trump_words$id))

```

```

[1] 1172

```

```
> # Put source back into the dataset
> (complete_with_source_tb =
+   inner_join(complete_tb, sources))
```

```
Joining, by = "id"
# A tibble: 8,790 x 5
  sentiment id          n source total
  <chr>     <chr>    <dbl> <chr> <int>
1 anger    676509769562251264 0. iPhone 3852
2 anger    680496083072593920 0. Android 4901
3 anger    680503951440121856 1. Android 4901
# ... with 8,787 more rows
```

```
> # counts by source and sentiment
> words_by_source_sentiment =
+   complete_with_source_tb %>%
+   group_by(source, sentiment, total) %>%
+   summarize(counts = sum(n)) %>%
+   ungroup()
>
> words_tidy_source_sentiment
```

- Everything together in a single chunk


```
> sources = trump_words %>%
+   group_by(source) %>%
+   mutate(total = n()) %>%
+   ungroup() %>%
+   distinct(id, source, total)
>
> words_by_source_sentiment = trump_words %>%
+   inner_join(nrc, by = "word") %>%
+   count(sentiment, id) %>%
+   ungroup() %>%
+   complete(sentiment, id, fill = list(n = 0)) %>%
+   inner_join(sources) %>%
+   group_by(source, sentiment, total) %>%
+   summarize(counts = sum(n)) %>%
+   ungroup()
>
> words_tidy_source_sentiment
```



```
# A tibble: 20 x 4
  source sentiment    total counts
  <chr>   <chr>      <int> <dbl>
1 Android anger        4901   321.
2 Android anticipation  4901   256.
3 Android disgust      4901   207.
4 Android fear          4901   268.
5 Android joy           4901   199.
6 Android negative     4901   560.
7 Android positive     4901   555.
8 Android sadness      4901   303.
9 Android surprise     4901   159.
10 Android trust        4901   369.
11 iPhone  anger        3852   169.
12 iPhone  anticipation  3852   172.
13 iPhone  disgust      3852    97.
14 iPhone  fear          3852   135.
15 iPhone  joy           3852   156.
16 iPhone  negative     3852   260.
17 iPhone  positive     3852   412.
18 iPhone  sadness      3852   147.
19 iPhone  surprise     3852   107.
20 iPhone  trust        3852   257.
```

```
> ggplot(words_by_source_sentiment, aes(
+   source, counts/total, fill = source)) +
+   geom_bar(stat = "identity", position = "dodge") +
+   labs(y = "Proportion", fill = "Source") +
+   facet_grid(~sentiment)
```

- It seems there is a clear difference in the category “negative.”




- However, graphical analysis alone is not enough for other categories.
- This is a count data, so let us test by assuming Poisson assumptions.

```
> sentiment_differences
+   words_by_source_sentiment %>%
+   group_by(sentiment) %>%
+   do(broom::tidy(poisson.test(.$counts, .$total)))
>
> sentiment_differences
```

```
# A tibble: 10 x 9
# Groups:   sentiment [10]
  sentiment estimate statistic p.value parameter conf.low conf.high method
  <chr>      <dbl>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <fct>
1 anger      1.49      321.  2.19e- 5    274.     1.24     1.81 Compari?
2 anticipat? 1.17      256.  1.19e- 1    240.     0.960    1.43 Compari?
3 disgust    1.68      207.  1.78e- 5    170.     1.31     2.16 Compari?
4 fear       1.56      268.  1.89e- 5    226.     1.26     1.93 Compari?
5 joy        1.00      199.  1.00e+ 0    199.     0.809    1.24 Compari?
6 negative   1.69      560.  7.09e-13    459.     1.46     1.97 Compari?
7 positive   1.06      555.  3.82e- 1    541.     0.930    1.21 Compari?
8 sadness    1.62      303.  1.15e- 6    252.     1.33     1.99 Compari?
9 surprise   1.17      159.  2.17e- 1    149.     0.908    1.51 Compari?
10 trust     1.13      369.  1.47e- 1    351.     0.960    1.33 Compari?
# ... with 1 more variable: alternative <fct>
```

- And we can visualise the difference with a 95% confidence interval:

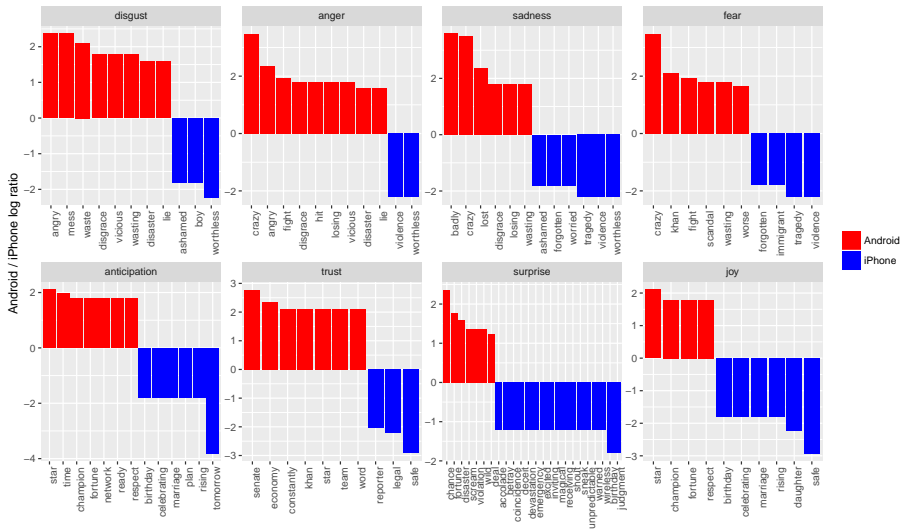


```

> sentiment_differences %>%
+   ungroup() %>%
+   mutate(sentiment =
+             reorder(sentiment, estimate)) %>%
+   mutate_at(c("estimate",
+               "conf.low",
+               "conf.high"),
+             funs(. - 1)) %>%
+   ggplot(aes(estimate, sentiment)) +
+   geom_point() +
+   geom_errorbarh(aes(
+     xmin = conf.low, xmax = conf.high)) +
+   scale_x_continuous(
+     labels = scales::percent_format()) +
+   labs(
+     x = "% increase in Android relative to iPhone",
+     y = "Sentiment")

```

Q: Which words in each category are driving those differences?



```

> android_iphone_ratios %>%
+   inner_join(nrc, by = "word") %>%
+   filter(!sentiment %in%
+           c("positive", "negative")) %>%
+   mutate(sentiment = reorder(sentiment, -logratio),
+           word = reorder(word, -logratio)) %>%
+   group_by(sentiment) %>%
+   top_n(10, abs(logratio)) %>% ungroup() %>%
+   ggplot(aes(
+     word, logratio, fill = logratio < 0)) +
+   facet_wrap(~ sentiment,
+             scales = "free", nrow = 2) +
+   geom_bar(stat = "identity") +
+   theme(axis.text.x =
+         element_text(angle = 90, hjust = 1)) +
+   labs(x = "",
+        y = "Android / iPhone log ratio") +
+   scale_fill_manual(
+     name = "", values = c("red", "blue"),
+     labels = c("Android", "iPhone"))


```


- beta's here are components of ϕ_k

```
> ap_lda = LDA(Trump, k = 2,  
+             method = "Gibbs",  
+             control = list(seed = 1234))  
>  
> ap_topics = tidy(ap_lda, matrix = "beta")  
> ap_topics
```

```
# A tibble: 5,068 x 3  
  topic term      beta  
  <int> <chr>    <dbl>  
1     1 american 0.0000219  
2     2 american 0.00343  
3     1 donald  0.00177  
4     2 donald  0.0000213  
5     1 great   0.0000219  
6     2 great   0.000235  
7     1 mass    0.0000219  
8     2 mass    0.000235  
9     1 s       0.0000219  
10    2 s       0.000874  
# with 5,058 more rows
```

- It is often interesting to see terms that are most common within each topic.





- gamma's here are components of θ_i

```
> ap_documents = tidy(ap_lda, matrix = "gamma")  
> ap_documents
```

```
# A tibble: 2,344 x 3  
  document      topic gamma  
  <chr>      <int> <dbl>  
1 759381869267980288      1 0.441  
2 762106904436961280      1 0.517  
3 758492727583576064      1 0.483  
4 746895602798178304      1 0.491  
5 697182075045179392      1 0.474  
6 701779181986680832      1 0.426  
7 747105049352888320      1 0.483  
8 750865499660091392      1 0.483  
9 761757988516401152      1 0.464  
10 754747397700485120      1 0.459  
# with 2,334 more rows
```

```
> ap_documents[  
+   ap_documents$document==759381869267980288,]
```

```
# A tibble: 2 x 3  
  document      topic gamma  
  <chr>      <int> <dbl>  
1 759381869267980288      1 0.441  
2 759381869267980288      2 0.559
```

```
> trump_source_tb[
+   trump_source_tb$id == 759381869267980288, 2]
```

```
# A tibble: 1 x 1
  source
  <chr>
1 iPhone
```

```
> trump_source_tb[
+   trump_source_tb$id == 759381869267980288, 3]
```

```
# A tibble: 1 x 1
  text
  <chr>
1 American homeownership rate in Q2 2016 was 62.9% - lowest rate in 51yrs. WE will br?
```

- It is not clear whether LDA has correctly sorted the document in this case.
- We need more content, or words in each document in order to

```
# A tibble: 2 x 3
  document      topic gamma
  <chr>         <int> <dbl>
1 759381869267980288      1 0.441
2 759381869267980288      2 0.559
```

have a decisive cut like the one we have seen in the news articles case.